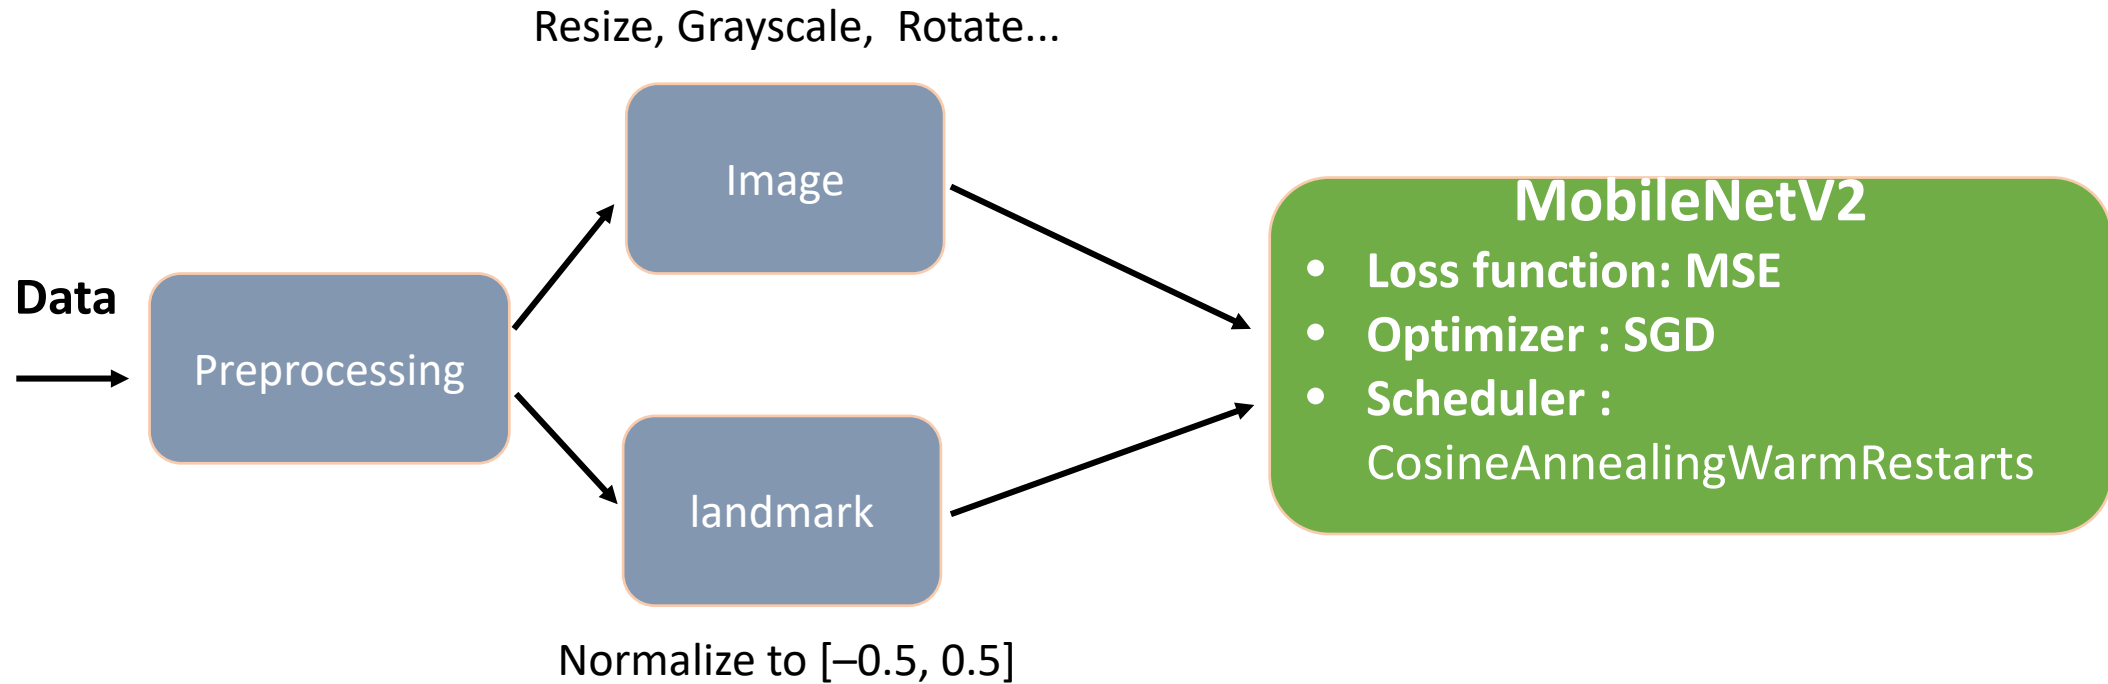


# Computer Vision Light-Weight Facial Landmark Prediction Challenge

Team AAA

蕭郁澄 r09942168, 楊欣睿 r10942079, 郭承賢 r09943098

# Global picture and results

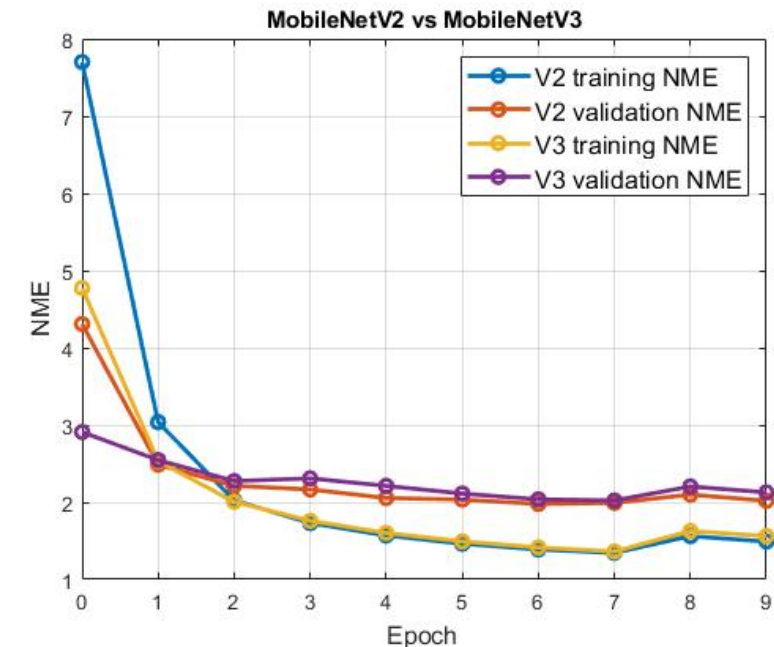


Leaderboard Rank	Validation NME	Public testing NME	Private testing NME
4	1.87095	2.03235	2.024754

# Model selection

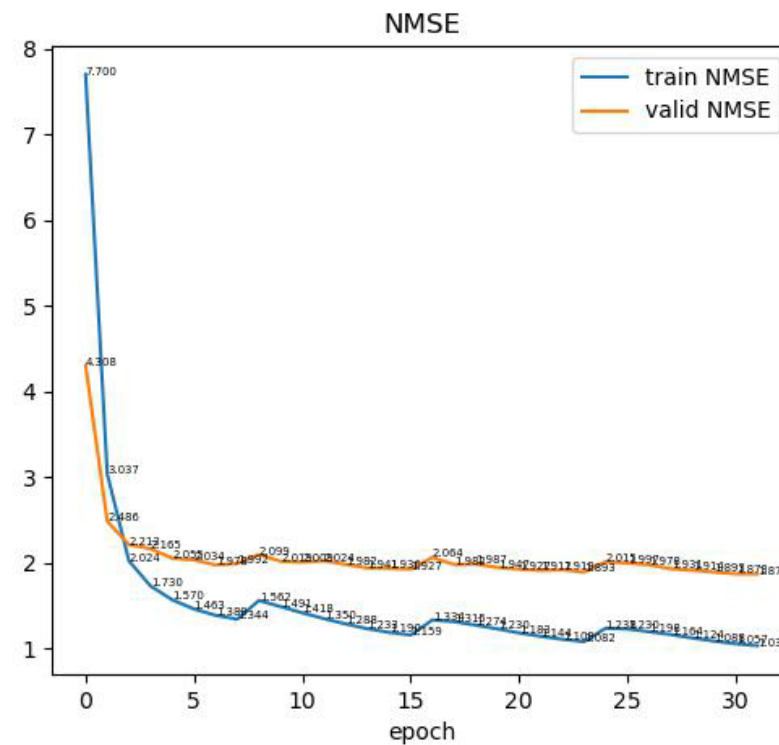

- It is well known that the MobileNet family performs well in the computer vision. Although [1] shows the MobileNetV3 outperforms the MobileNetV2 on the image classification tasks.
- After the exhaustive experiments, we found the **MobileNetV3 is easier to overfit** the training dataset and the **MobileNetV2 is more suitable** for our task.
- Due to the model size constraint(15MB), we choose the **1.25 multiplier MobileNetV2 (14.25 MB)** as our training model.

Network	Top-1	MAdds	Params	P-1	P-2	P-3
V3-Large 1.0	<b>75.2</b>	<b>219</b>	5.4M	<b>51</b>	<b>61</b>	<b>44</b>
V3-Large 0.75	73.3	155	4.0M	39	46	40
MnasNet-A1	75.2	315	3.9M	71	86	61
Proxyless[5]	74.6	320	4.0M	72	84	60
V2 1.0	72.0	300	3.4M	64	76	56
V3-Small 1.0	<b>67.4</b>	<b>56</b>	2.5M	<b>15.8</b>	<b>19.4</b>	<b>14.4</b>
V3-Small 0.75	65.4	44	2.0M	12.8	15.6	11.7
Mnas-small [43]	64.9	65.1	1.9M	20.3	24.2	17.2
V2 0.35	60.8	59.2	1.6M	16.6	19.6	13.9



# Scheduler

- `optim.lr_scheduler.CosineAnnealingWarmRestarts(optimizer, T_0 = 8, T_mult = 1)`



# Preprocessing

- We adopt the following preprocessing for landmarks and images, respectively.

- Landmarks:

- Normalized to  $[-0.5, 0.5]$

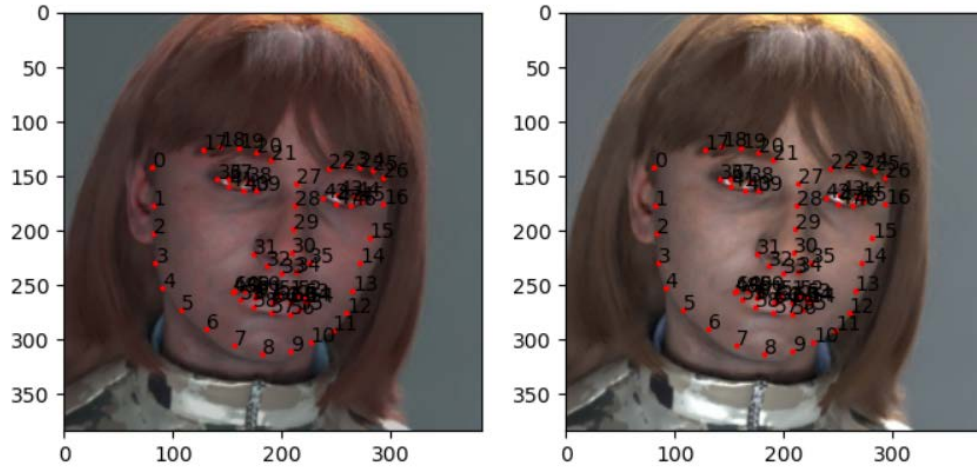


- Image:

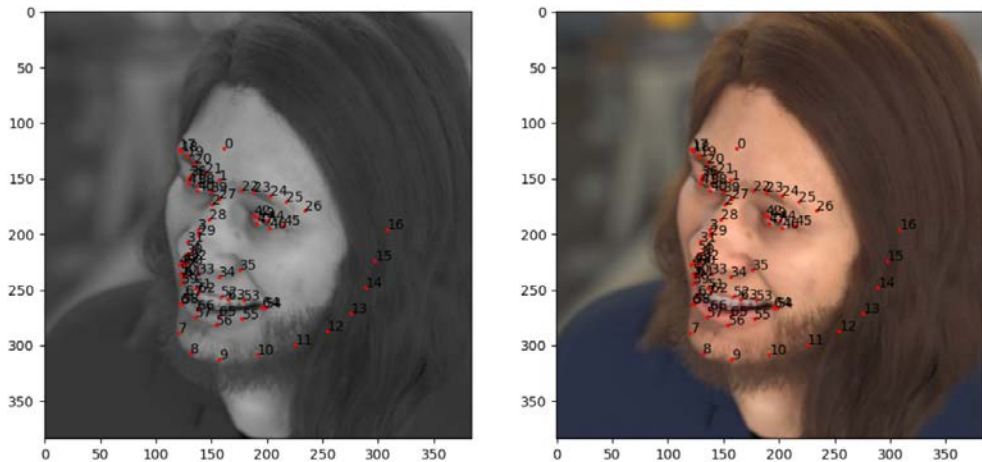
- RandColorJitter, RandGrayscale, Rrandom\_horizontalflip
  - RandomRotation, Random\_scale, Random\_erasemargin, RandomBlur, RandomErasing

# Visualize the preprocessing

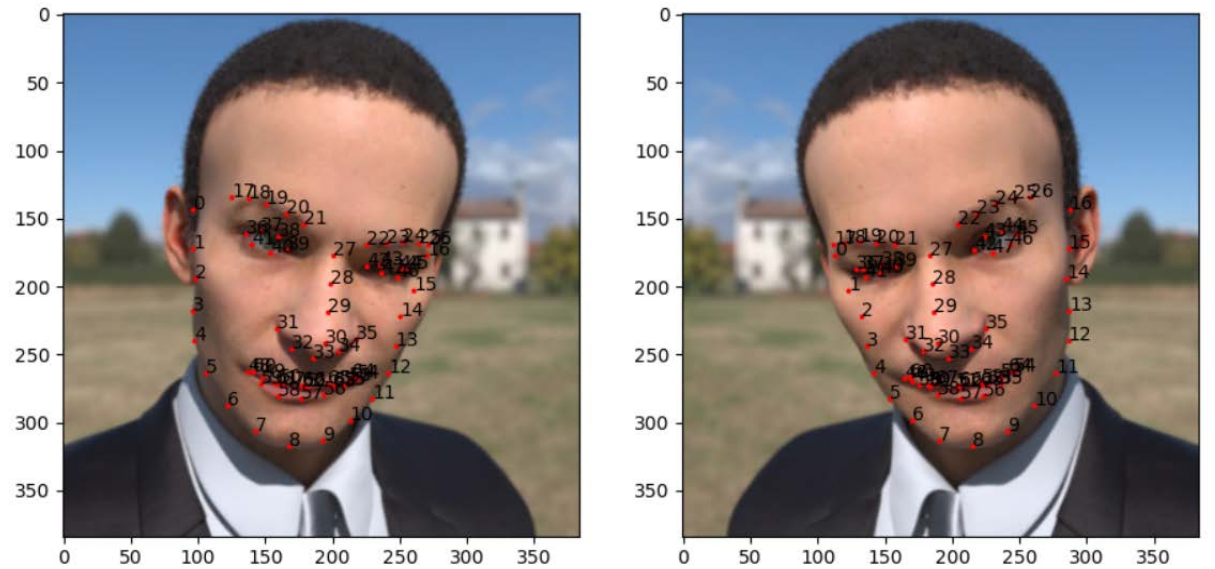
- **ColorJitter**



- **RandomGrayscale**



- **Random\_horizontalflip**




# Difference of dataset

- Train on synthetic data, test on real image

Tilted head



Black margin



Low resolution



Blocked by something




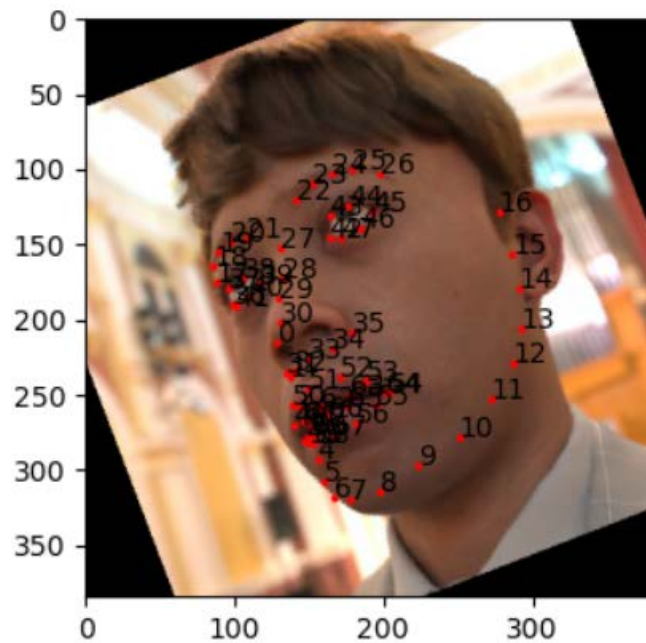


# Difference of dataset

Tilted head




- **Random\_rotate**



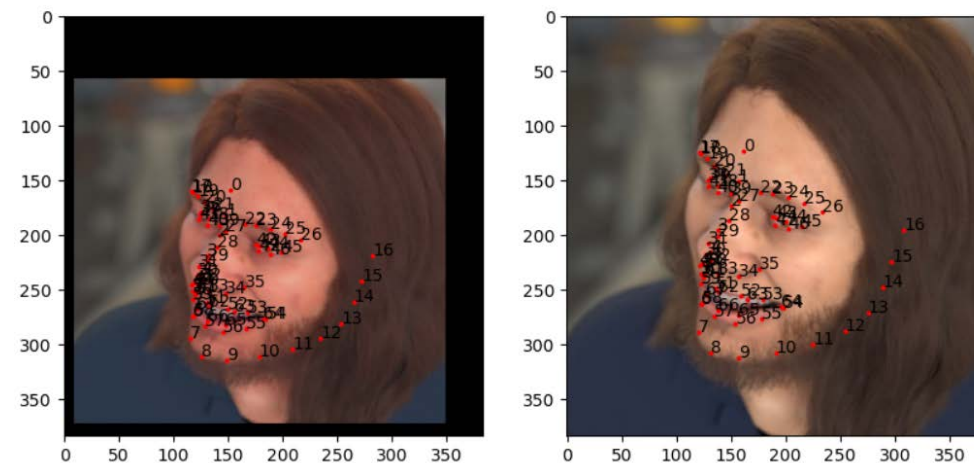


# Difference of dataset


Black margin



- **Random\_scale**



- **Random\_erasemargin**

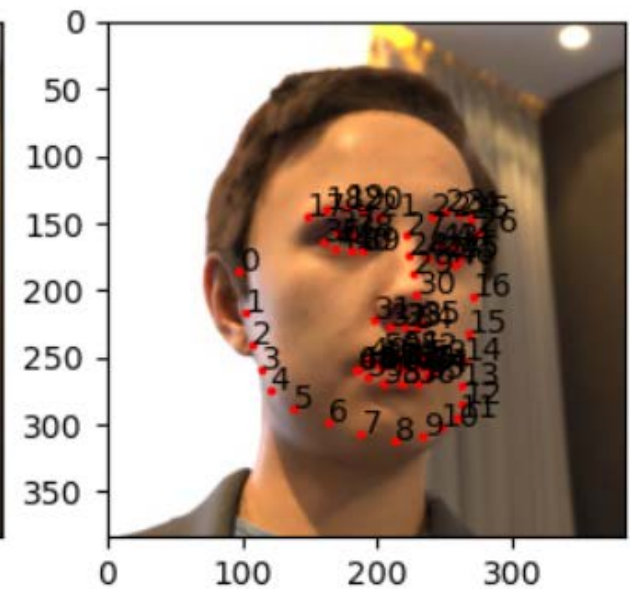



# Difference of dataset

Low resolution



- **Random\_blur**


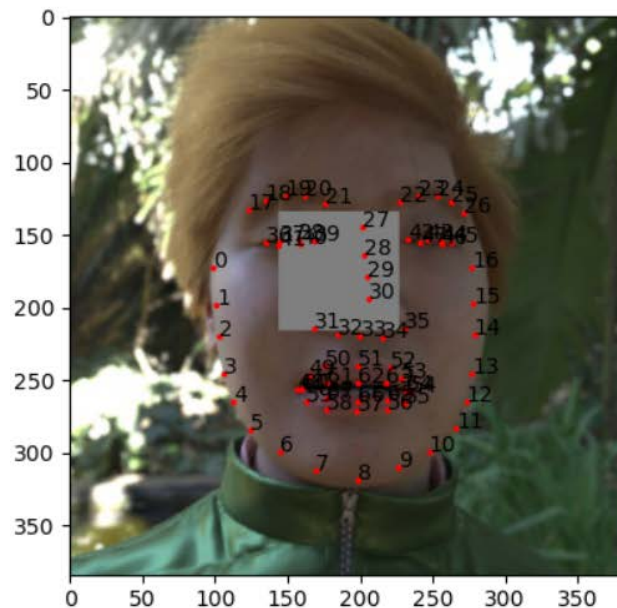


# Difference of dataset

Blocked by something




- **RandomErasing**



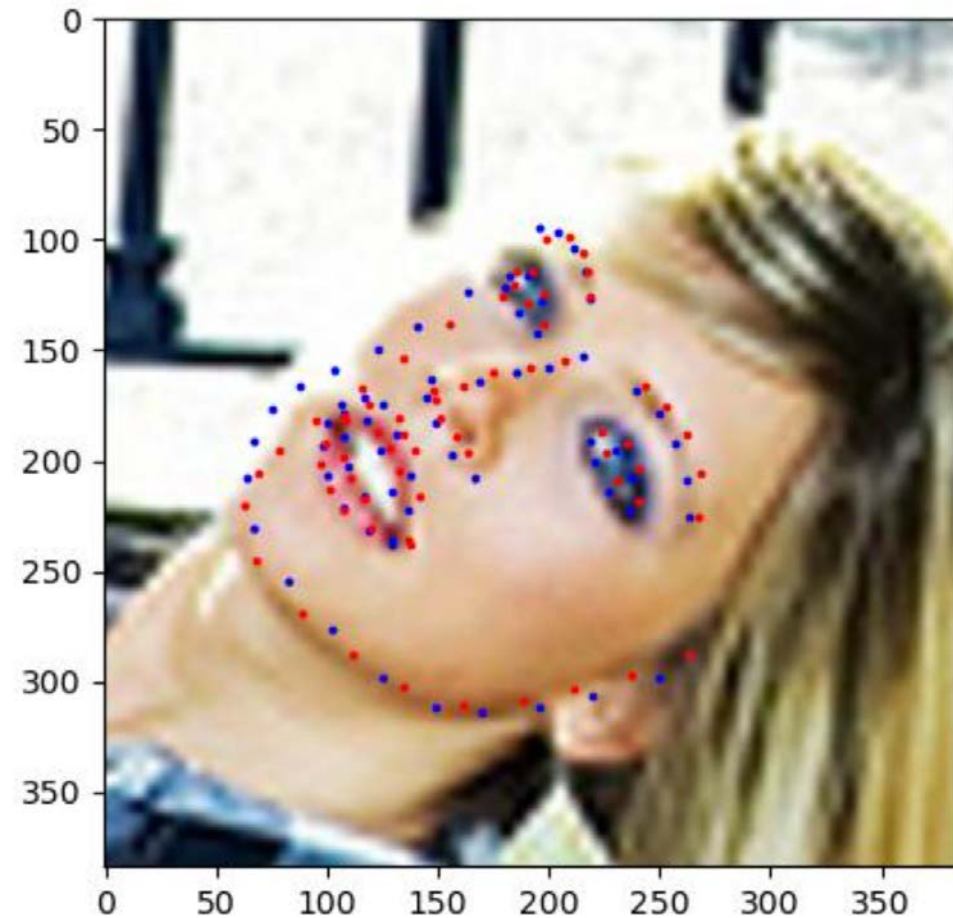


# Data augmentation (red : ground truth, blue: predicted landmarks)

- Without data augmentation

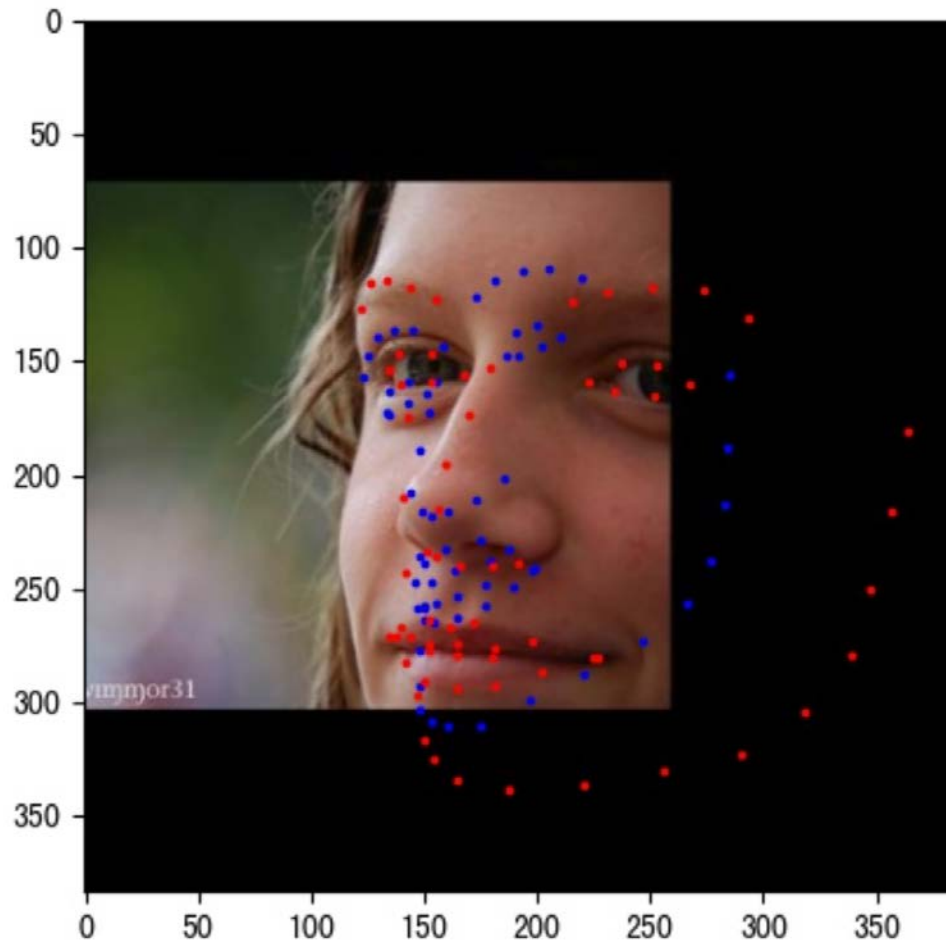


- With data augmentation

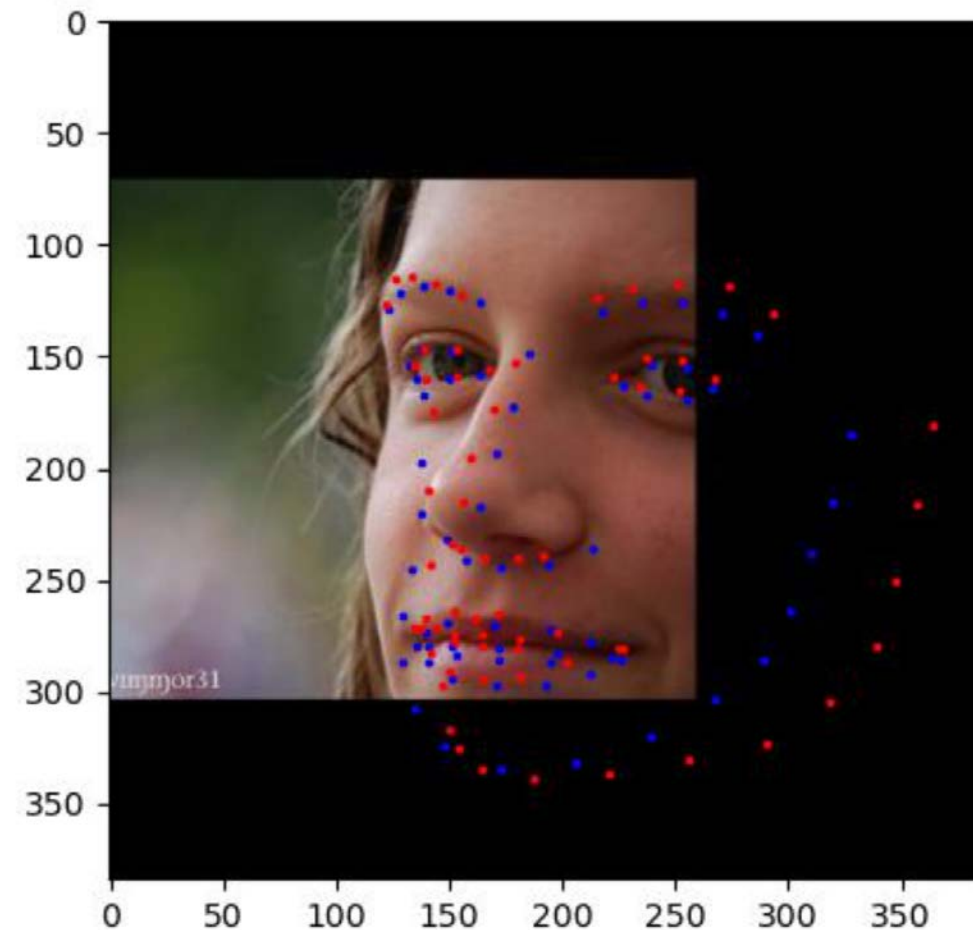


# Data augmentation (red : ground truth, blue: predicted landmarks)

- Without data augmentation




- With data augmentation



# Experiment : resize 384x384 image to 224x224

- Pros : reduce the training time by half and it's convenient to tune hyperparameters.
- Cons : loss of precision.

Except the image size, other parameters are same.







Thank you

# Appendix: Fine-tune Hyper prameter

```
hyper = {  
    'batch_size': 16,  
    'img_resize': 384,  
    'lr': 0.005,  
    'num_epoch': 32,  
    'cos_T_0': 8,  
    'cos_T_multi': 1,  
}
```

```
p_dict = {  
    'ColorJitter_b': 0.3,  
    'ColorJitter_c': 0.3,  
    'ColorJitter_s': 0.3,  
    'ColorJitter_h': 0.1,  
    'RandomGrayscale': 0.4,  
    'random_scale_p': 0,  
    'random_scale_scale': 0.8,  
    'random_rotate_p': 0.4,  
    'random_rotate_degree': 45,  
    'random_blur_p': 0.2,  
    'random_blur_min': 128,  
    'random_horizontalflip_p': 0.5,  
    'random_erasemargin_p': 0.4,  
    'random_erasemargin_ratio': 0.7,  
    'RandomErasing_p': 0.1,  
    'RandomErasing_scale_min': 0.02,  
    'RandomErasing_scale_max': 0.1,  
}
```